

Table 1: On-device inference efficiency and accuracy on Raspberry Pi 4 (Cortex-A72 @ 1.5 GHz). Best metrics are **bold**, second best are underlined. Lower is better ($\downarrow$) for computational cost, latency, memory footprint, and energy.

Model	Top-1 Acc	FLOPs $\downarrow$	Latency $\downarrow$	Peak RAM $\downarrow$	Energy $\downarrow$
MobileNetV3-Large	75.2%	219M	8.4 ms	42.1 MB	18.2 mJ
EfficientNet-Lite0	75.1%	385M	11.2 ms	58.4 MB	24.6 mJ
ShuffleNetV2-1.5x	72.6%	146M	6.1 ms	36.8 MB	13.5 mJ
GhostNetV2-1.3x	<u>75.3%</u>	167M	<u>7.2 ms</u>	<u>39.5 MB</u>	<u>15.1 mJ</u>
EdgeNeXt-Small	79.4%	260M	10.8 ms	51.2 MB	22.8 mJ