

Table 1: Robustness benchmark under ImageNet-C corruptions. Accuracy columns (Clean, Noise, Blur, Weather) are shaded with independent per-column *Blues* heatmaps. Arrows indicate direction of improvement ($\uparrow$ higher is better, $\downarrow$ lower is better). Lowest Mean Corruption Error (mCE $\downarrow$) is **bold**, second lowest is underlined.

Model	Params $\downarrow$	Clean (%) $\uparrow$	Noise (%) $\uparrow$	Blur (%) $\uparrow$	Weather (%) $\uparrow$	mCE $\downarrow$
ResNet-50	25.6M	76.1	42.1	48.4	54.2	76.8
ConvNeXt-T	28.6M	82.1	55.4	60.1	66.8	<u>62.4</u>
DeiT-S	22.1M	79.8	51.2	54.6	61.3	69.1
Swin-T	28.3M	81.2	57.8	62.4	68.1	59.8
EfficientNet-B3	12.2M	81.6	52.3	57.8	63.4	65.2