Table 1: Natural language understanding benchmark results on GLUE tasks (reported as mean ± SEM over 5 seeds). Best results within each model group are **bold**, second best are <u>underlined</u>. Arrows indicate direction of improvement (↑ higher is better, ↓ lower is better). The top overall model in each category is highlighted in light gray.

Model	**MNLI ↑**	**QNLI ↑**	**SST-2 ↑**	**CoLA ↑**	**Average ↑**	**Latency ↓**
Supervised Fine-Tuning						
BERT-Base	84.6 ± 0.7	90.8 ± 0.5	92.3 ± 0.4	58.4 ± 1.1	81.5 ± 0.5	**14.8 ms**
RoBERTa-Base	<u>87.8 ± 0.5</u>	<u>92.7 ± 0.4</u>	<u>94.8 ± 0.4</u>	<u>63.6 ± 0.9</u>	<u>84.7 ± 0.4</u>	<u>15.2 ms</u>
DeBERTa-v3	**89.9 ± 0.4**	**94.1 ± 0.3**	**95.6 ± 0.3**	**69.5 ± 0.7**	**87.3 ± 0.3**	17.4 ms
Unsupervised Representations						
SimCSE-BERT	81.2 ± 0.8	88.5 ± 0.6	90.1 ± 0.5	52.1 ± 1.3	78.0 ± 0.6	**14.5 ms**
ELECTRA-Base	<u>88.5 ± 0.4</u>	<u>93.2 ± 0.4</u>	<u>95.1 ± 0.3</u>	<u>67.2 ± 0.8</u>	<u>86.0 ± 0.4</u>	<u>15.0 ms</u>
LLaMA-Distill	**91.2 ± 0.4**	**94.8 ± 0.3**	**96.2 ± 0.2**	**71.8 ± 0.6**	**88.5 ± 0.3**	21.5 ms