

Table 1: Natural language understanding benchmark results on GLUE tasks (reported as mean $\pm$ SEM over 5 seeds). Best results are **bold**, second best are underlined. The top overall average model is highlighted in light gray.

Model	MNLI	QNLI	SST-2	CoLA	Average	Latency
BERT-Base	84.6 $\pm$ 0.7	90.8 $\pm$ 0.5	92.3 $\pm$ 0.4	58.4 $\pm$ 1.1	81.5 $\pm$ 0.5	14.8 ms
RoBERTa-Base	87.8 $\pm$ 0.5	92.7 $\pm$ 0.4	94.8 $\pm$ 0.4	63.6 $\pm$ 0.9	84.7 $\pm$ 0.4	15.2 ms
ELECTRA-Base	88.5 $\pm$ 0.4	93.2 $\pm$ 0.4	95.1 $\pm$ 0.3	67.2 $\pm$ 0.8	86.0 $\pm$ 0.4	<u>15.0 ms</u>
DeBERTa-v3	<u>89.9 $\pm$ 0.4</u>	<u>94.1 $\pm$ 0.3</u>	<u>95.6 $\pm$ 0.3</u>	<u>69.5 $\pm$ 0.7</u>	<u>87.3 $\pm$ 0.3</u>	17.4 ms
LLaMA-1B-Distill	91.2 $\pm$ 0.4	94.8 $\pm$ 0.3	96.2 $\pm$ 0.2	71.8 $\pm$ 0.6	88.5 $\pm$ 0.3	21.5 ms